

An introduction to tidycensus

Ram Mandava

2023-02-17

1. What worked well for you and what didn't work?

Analyzed the US Census Data with help of existing R Studio code chunks and produce in terms of plots, tables, and caching data. These complex tasks can be segregated into sections to produce the results and converge in a publication (along with supplementary materials). Final “enriched document”, can back-trace and reproduce the entire analysis made or just part of it.

If we have a good understanding of R programming, the maintainability and **reproducibility** of outcomes will be much easier to handle.

But some of the R code chunks didn't work and gave the below errors and were unable to fix the issues.

Examples:

1. The API message returned is the error: unknown variable 'P01001'.
2. Error in Use Method("gather"): no applicable method for 'gather' applied to an object of class "character".

2. What do you see as the value of using RMarkdown?

R Markdown is a subset of literate programming, which aims to make it simple to write reproducible web-based reports by combining text and source code into a single document. Used to create documents in many different formats, including HTML, PDF, and MS Word.

Ability to break down into sections that give better readability. Both papers and supplementary materials can be enriched by documents that keep track of all code, data, and results produced during the analysis.

3. Based on lab 2 so far, what things would you like to do?

Familiarized with the United States Census and the US Census Bureau, and how Census data can be accessed and used by analysts.

Importance of R package for working with US Census Bureau data in a tidy format.

Understanding the basic data requests with the package and various options in the package.

R programming techniques to streamline the data wrangling process for spatial Census data analysis and explore US Census data with visualization. Workflows and best practices for preparing data for visualization and building charts for presenting Census data analyses.

Getting started with tidycensus code chunks with results

```
library(tidycensus)
```

```
## Warning: package 'tidycensus' was built under R version 4.2.2
```

```
#census_api_key("1840a214a15b3834c5308103bdaf17009f4db1b4", install = TRUE)
```

```
total_population_10 <- get_decennial(  
  geography = "state",  
  variables = "P001001",  
  year = 2010  
)
```

```
## Getting data from the 2010 decennial Census
```

```
## Using Census Summary File 1
```

```
print(total_population_10)
```

```
## # A tibble: 52 × 4  
##   GEOID NAME      variable    value  
##   <chr> <chr>      <chr>      <dbl>  
## 1 01    Alabama    P001001    4779736  
## 2 02    Alaska     P001001     710231  
## 3 04    Arizona    P001001    6392017  
## 4 05    Arkansas   P001001    2915918  
## 5 06    California P001001    37253956  
## 6 22    Louisiana  P001001    4533372  
## 7 21    Kentucky   P001001    4339367  
## 8 08    Colorado   P001001    5029196  
## 9 09    Connecticut P001001    3574097  
## 10 10   Delaware   P001001     897934  
## # ... with 42 more rows
```

```
aian_2020 <- get_decennial(  
  geography = "state",  
  variables = "P1_005N",  
  year = 2020,  
  sumfile = "p1"  
)
```

```
## Getting data from the 2020 decennial Census
```

```
## Using the PL 94-171 Redistricting Data summary file
```

```
## Note: 2020 decennial Census data use differential privacy, a technique that
## introduces errors into data to preserve respondent confidentiality.
## i Small counts should be interpreted with caution.
## i See https://www.census.gov/library/fact-sheets/2021/protecting-the-confidentiality-of-the-2020-census-redistricting-data.html for additional guidance.
## This message is displayed once per session.
```

```
print(aian_2020)
```

```
## # A tibble: 52 × 4
##   GEOID NAME          variable value
##   <chr> <chr>          <chr>   <dbl>
## 1 42    Pennsylvania    P1_005N  31052
## 2 06    California      P1_005N 631016
## 3 54    West Virginia    P1_005N   3706
## 4 49    Utah             P1_005N  41644
## 5 36    New York         P1_005N 149690
## 6 11    District of Columbia P1_005N   3193
## 7 02    Alaska           P1_005N 111575
## 8 12    Florida          P1_005N  94795
## 9 45    South Carolina    P1_005N  24303
## 10 38   North Dakota      P1_005N  38914
## # ... with 42 more rows
```

```
born_in_mexico <- get_acs(
  geography = "state",
  variables = "B05006_150",
  year = 2020
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
print(born_in_mexico)
```

```
## # A tibble: 52 × 5
##   GEOID NAME          variable estimate moe
##   <chr> <chr>          <chr>      <dbl> <dbl>
## 1 01    Alabama      B05006_150    46927  1846
## 2 02    Alaska       B05006_150     4181   709
## 3 04    Arizona      B05006_150   510639  8028
## 4 05    Arkansas      B05006_150    60236  2182
## 5 06    California    B05006_150  3962910 25353
## 6 08    Colorado      B05006_150   215778  4888
## 7 09    Connecticut    B05006_150    28086  2144
## 8 10    Delaware      B05006_150    14616  1065
## 9 11    District of Columbia B05006_150     4026   761
## 10 12   Florida      B05006_150   257933  6418
## # ... with 42 more rows
```

```
born_in_mexico_1yr <- get_acs(
  geography = "state",
  variables = "B05006_150",
  survey = "acs1",
  year = 2019
)
```

```
## Getting data from the 2019 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
print(born_in_mexico_1yr)
```

```
## # A tibble: 52 × 5
##   GEOID NAME          variable estimate moe
##   <chr> <chr>          <chr>      <dbl> <dbl>
## 1 01    Alabama      B05006_150      NA    NA
## 2 02    Alaska       B05006_150      NA    NA
## 3 04    Arizona      B05006_150   516618 15863
## 4 05    Arkansas      B05006_150      NA    NA
## 5 06    California    B05006_150  3951224 40506
## 6 08    Colorado      B05006_150   209408 12214
## 7 09    Connecticut    B05006_150    26371  4816
## 8 10    Delaware      B05006_150      NA    NA
## 9 11    District of Columbia B05006_150      NA    NA
## 10 12   Florida      B05006_150   261614 17571
## # ... with 42 more rows
```

```
age_table <- get_acs(  
  geography = "state",  
  table = "B01001",  
  year = 2020  
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
## Loading ACS5 variables for 2020 from table B01001. To cache this dataset for faster access to  
ACS tables in the future, run this function with `cache_table = TRUE`. You only need to do this  
once per ACS dataset.
```

```
print(age_table)
```

```
## # A tibble: 2,548 × 5  
##   GEOID NAME    variable  estimate  moe  
##   <chr> <chr>    <chr>      <dbl> <dbl>  
## 1 01    Alabama B01001_001 4893186    NA  
## 2 01    Alabama B01001_002 2365734  1090  
## 3 01    Alabama B01001_003 149579   672  
## 4 01    Alabama B01001_004 150937  2202  
## 5 01    Alabama B01001_005 160287  2159  
## 6 01    Alabama B01001_006 96832   565  
## 7 01    Alabama B01001_007 65459   961  
## 8 01    Alabama B01001_008 36705  1467  
## 9 01    Alabama B01001_009 33089  1547  
## 10 01   Alabama B01001_010 93871  2045  
## # ... with 2,538 more rows
```

```
cbsa_population <- get_acs(  
  geography = "cbsa",  
  variables = "B01003_001",  
  year = 2020  
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
## Getting data from the 2016-2020 5-year ACS  
print(cbsa_population)
```

```
## # A tibble: 939 × 5
##   GEOID NAME                variable estimate moe
##   <chr> <chr>                <chr>      <dbl> <dbl>
## 1 10100 Aberdeen, SD Micro Area B01003_001 42864 NA
## 2 10140 Aberdeen, WA Micro Area B01003_001 73769 NA
## 3 10180 Abilene, TX Metro Area  B01003_001 171354 NA
## 4 10220 Ada, OK Micro Area      B01003_001 38385 NA
## 5 10300 Adrian, MI Micro Area   B01003_001 98310 NA
## 6 10380 Aguadilla-Isabela, PR Metro Area B01003_001 295172 NA
## 7 10420 Akron, OH Metro Area    B01003_001 703286 NA
## 8 10460 Alamogordo, NM Micro Area B01003_001 66804 NA
## 9 10500 Albany, GA Metro Area    B01003_001 147431 NA
## 10 10540 Albany-Lebanon, OR Metro Area B01003_001 127216 NA
## # ... with 929 more rows
```

```
wi_income <- get_acs(
  geography = "county",
  variables = "B19013_001",
  state = "WI",
  year = 2020
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
print(wi_income)
```

```
## # A tibble: 72 × 5
##   GEOID NAME                variable estimate moe
##   <chr> <chr>                <chr>      <dbl> <dbl>
## 1 55001 Adams County, Wisconsin B19013_001 48906 2387
## 2 55003 Ashland County, Wisconsin B19013_001 47869 3190
## 3 55005 Barron County, Wisconsin B19013_001 52346 2092
## 4 55007 Bayfield County, Wisconsin B19013_001 57257 2496
## 5 55009 Brown County, Wisconsin  B19013_001 64728 1419
## 6 55011 Buffalo County, Wisconsin B19013_001 58364 1871
## 7 55013 Burnett County, Wisconsin B19013_001 53555 2513
## 8 55015 Calumet County, Wisconsin B19013_001 76065 2314
## 9 55017 Chippewa County, Wisconsin B19013_001 61215 2064
## 10 55019 Clark County, Wisconsin  B19013_001 54463 1089
## # ... with 62 more rows
```

```
dane_income <- get_acs(
  geography = "tract",
  variables = "B19013_001",
  state = "WI",
  county = "Dane",
  year = 2020
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
print(dane_income)
```

```
## # A tibble: 125 × 5
##   GEOID      NAME      variable estim...1 moe
##   <chr>      <chr>      <chr>      <dbl> <dbl>
## 1 55025000100 Census Tract 1, Dane County, Wisconsin B19013_0... 74054 15662
## 2 55025000201 Census Tract 2.01, Dane County, Wisconsin B19013_0... 92460 27067
## 3 55025000202 Census Tract 2.02, Dane County, Wisconsin B19013_0... 88092 5189
## 4 55025000204 Census Tract 2.04, Dane County, Wisconsin B19013_0... 82717 12175
## 5 55025000205 Census Tract 2.05, Dane County, Wisconsin B19013_0... 100000 17506
## 6 55025000301 Census Tract 3.01, Dane County, Wisconsin B19013_0... 37016 11524
## 7 55025000302 Census Tract 3.02, Dane County, Wisconsin B19013_0... 117321 28723
## 8 55025000401 Census Tract 4.01, Dane County, Wisconsin B19013_0... 100434 12108
## 9 55025000402 Census Tract 4.02, Dane County, Wisconsin B19013_0... 105850 12205
## 10 55025000406 Census Tract 4.06, Dane County, Wisconsin B19013_0... 74009 2811
## # ... with 115 more rows, and abbreviated variable name 1estimate
```

```
nrow(wi_income)
```

```
## [1] 72
```

```
wi_income_1yr <- get_acs(
  geography = "county",
  variables = "B19013_001",
  state = "WI",
  year = 2019,
  survey = "acs1"
)
```

```
## Getting data from the 2019 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
print(wi_income_1yr)
```

```
## # A tibble: 23 × 5
##   GEOID NAME                variable estimate moe
##   <chr> <chr>                <chr>      <dbl> <dbl>
## 1 55009 Brown County, Wisconsin B19013_001 64458 3713
## 2 55025 Dane County, Wisconsin B19013_001 77504 3180
## 3 55027 Dodge County, Wisconsin B19013_001 61300 2119
## 4 55035 Eau Claire County, Wisconsin B19013_001 65662 3914
## 5 55039 Fond du Lac County, Wisconsin B19013_001 65329 4646
## 6 55055 Jefferson County, Wisconsin B19013_001 71708 5901
## 7 55059 Kenosha County, Wisconsin B19013_001 65997 4897
## 8 55063 La Crosse County, Wisconsin B19013_001 58921 3969
## 9 55071 Manitowoc County, Wisconsin B19013_001 60785 3913
## 10 55073 Marathon County, Wisconsin B19013_001 65851 2868
## # ... with 13 more rows
```

```
nrow(wi_income_1yr)
```

```
## [1] 23
```

```
v16 <- load_variables(2016, "acs5", cache = TRUE)

#View(v16)
```

```
hhinc <- get_acs(
  geography = "state",
  table = "B19001",
  survey = "acs1",
  year = 2016
)
```

```
## Getting data from the 2016 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Loading ACS1 variables for 2016 from table B19001. To cache this dataset for faster access to
## ACS tables in the future, run this function with `cache_table = TRUE`. You only need to do this
## once per ACS dataset.
```

```
print(hhinc)
```



```
## # A tibble: 884 × 5
##   GEOID NAME    variable estimate moe
##   <chr> <chr>    <chr>      <dbl> <dbl>
## 1 01 Alabama B19001_001 1852518 12189
## 2 01 Alabama B19001_002 176641 6328
## 3 01 Alabama B19001_003 120590 5347
## 4 01 Alabama B19001_004 117332 5956
## 5 01 Alabama B19001_005 108912 5308
## 6 01 Alabama B19001_006 102080 4740
## 7 01 Alabama B19001_007 103366 5246
## 8 01 Alabama B19001_008 91011 4699
## 9 01 Alabama B19001_009 86996 4418
## 10 01 Alabama B19001_010 74864 4210
## # ... with 874 more rows
```

```
hhinc_wide <- get_acs(
  geography = "state",
  table = "B19001",
  survey = "acs1",
  year = 2016,
  output = "wide"
)
```

```
## Getting data from the 2016 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Loading ACS1 variables for 2016 from table B19001. To cache this dataset for faster access to
ACS tables in the future, run this function with `cache_table = TRUE`. You only need to do this
once per ACS dataset.
```

```
print(hhinc_wide)
```

```
## # A tibble: 52 × 36
##   GEOID NAME      B1900...1 B1900...2 B1900...3 B1900...4 B1900...5 B1900...6 B1900...7 B1900...8
##   <chr> <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 28 Missis... 1091245    8803    113124    4835    87136    5004    71206    4058
## 2 29 Missou... 2372190   10844   160615    6705   122649    4654   123789    5201
## 3 30 Montana 416125    4426    26734    2183    24786    2391    22330    2391
## 4 31 Nebras... 747562    4452    45794    3116    33266    2466    31084    2533
## 5 32 Nevada 1055158    6433    68507    4886    42720    3071    53143    3653
## 6 33 New Ha... 520643    5191    20890    2566    15933    1908    18190    2315
## 7 34 New Je... 3194519   10274   170029    6836   118862    5855   123335    6065
## 8 35 New Me... 758364    6296    66983    4439    48930    3220    50025    4091
## 9 36 New Yo... 7209054   17665   543763   12132   352029    9607   322683    7756
## 10 37 North ... 3882423   16063   282491    7816   228088    7916   209825    6844
## # ... with 42 more rows, 26 more variables: B19001_005E <dbl>, B19001_005M <dbl>,
## # B19001_006E <dbl>, B19001_006M <dbl>, B19001_007E <dbl>, B19001_007M <dbl>,
## # B19001_008E <dbl>, B19001_008M <dbl>, B19001_009E <dbl>, B19001_009M <dbl>,
## # B19001_010E <dbl>, B19001_010M <dbl>, B19001_011E <dbl>, B19001_011M <dbl>,
## # B19001_012E <dbl>, B19001_012M <dbl>, B19001_013E <dbl>, B19001_013M <dbl>,
## # B19001_014E <dbl>, B19001_014M <dbl>, B19001_015E <dbl>, B19001_015M <dbl>,
## # B19001_016E <dbl>, B19001_016M <dbl>, B19001_017E <dbl>, ...
```

```
cimarron_blocks <- get_decennial(
  geography = "block",
  variables = "H1_001N",
  state = "OK",
  county = "Cimarron",
  year = 2020,
  sumfile = "pl"
)
```

```
## Getting data from the 2020 decennial Census
```

```
## Using the PL 94-171 Redistricting Data summary file
```

```
print(cimarron_blocks)
```

```
## # A tibble: 1,269 × 4
##   GEOID      NAME                                varia...1 value
##   <chr>      <chr>                                <chr>    <dbl>
## 1 400259501001984 Block 1984, Block Group 1, Census Tract 9501, ... H1_001N    0
## 2 400259503001035 Block 1035, Block Group 1, Census Tract 9503, ... H1_001N    0
## 3 400259503001068 Block 1068, Block Group 1, Census Tract 9503, ... H1_001N    5
## 4 400259503001146 Block 1146, Block Group 1, Census Tract 9503, ... H1_001N    0
## 5 400259503001197 Block 1197, Block Group 1, Census Tract 9503, ... H1_001N    7
## 6 400259503001218 Block 1218, Block Group 1, Census Tract 9503, ... H1_001N    2
## 7 400259501001067 Block 1067, Block Group 1, Census Tract 9501, ... H1_001N    0
## 8 400259501001118 Block 1118, Block Group 1, Census Tract 9501, ... H1_001N    0
## 9 400259501001141 Block 1141, Block Group 1, Census Tract 9501, ... H1_001N    0
## 10 400259501001223 Block 1223, Block Group 1, Census Tract 9501, ... H1_001N    0
## # ... with 1,259 more rows, and abbreviated variable name 1variable
```

```
ga <- get_acs(
  geography = "county",
  state = "Georgia",
  variables = c(medinc = "B19013_001",
               medage = "B01002_001"),
  year = 2020
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
print(ga)
```

```
## # A tibble: 318 × 5
##   GEOID NAME                                variable estimate    moe
##   <chr> <chr>                                <chr>    <dbl> <dbl>
## 1 13001 Appling County, Georgia medage      39.9    1.7
## 2 13001 Appling County, Georgia medinc     37924   4761
## 3 13003 Atkinson County, Georgia medage      35.9    1.5
## 4 13003 Atkinson County, Georgia medinc     35703   5493
## 5 13005 Bacon County, Georgia medage      36.5     1
## 6 13005 Bacon County, Georgia medinc     36692   3774
## 7 13007 Baker County, Georgia medage      52.2    4.8
## 8 13007 Baker County, Georgia medinc     34034   9879
## 9 13009 Baldwin County, Georgia medage      35.8    0.5
## 10 13009 Baldwin County, Georgia medinc     46250   4707
## # ... with 308 more rows
```

```
ga_wide <- get_acs(
  geography = "county",
  state = "Georgia",
  variables = c(medinc = "B19013_001",
                medage = "B01002_001"),
  output = "wide",
  year = 2020
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
print(ga_wide)
```

```
## # A tibble: 159 × 6
##   GEOID NAME                medincE medincM medageE medageM
##   <chr> <chr>                <dbl>   <dbl>   <dbl>   <dbl>
## 1 13001 Appling County, Georgia 37924   4761    39.9    1.7
## 2 13003 Atkinson County, Georgia 35703   5493    35.9    1.5
## 3 13005 Bacon County, Georgia 36692   3774    36.5     1
## 4 13007 Baker County, Georgia 34034   9879    52.2    4.8
## 5 13011 Banks County, Georgia 50912   4278    41.5    1.1
## 6 13013 Barrow County, Georgia 62990   2562     36     0.3
## 7 13017 Ben Hill County, Georgia 32077   4008    39.5    1.4
## 8 13021 Bibb County, Georgia 41317   1220    36.3    0.3
## 9 13023 Bleckley County, Georgia 46992   6279     36     1.5
## 10 13027 Brooks County, Georgia 37516   4438    43.6    0.9
## # ... with 149 more rows
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.2.2
```

```
## — Attaching packages — tidyverse 1.3.2 —
## ✓ ggplot2 3.4.0      ✓ purrr 1.0.1
## ✓ tibble 3.1.8       ✓ dplyr 1.0.10
## ✓ tidyr 1.3.0        ✓ stringr 1.5.0
## ✓ readr 2.1.3        ✓ forcats 1.0.0
```

```
## Warning: package 'ggplot2' was built under R version 4.2.2
```

```
## Warning: package 'tidyr' was built under R version 4.2.2
```

```
## Warning: package 'readr' was built under R version 4.2.2
```

```
## Warning: package 'purrr' was built under R version 4.2.2
```

```
## Warning: package 'stringr' was built under R version 4.2.2
```

```
## Warning: package 'forcats' was built under R version 4.2.2
```

```
## — Conflicts ————— tidyverse_conflicts() —  
## X dplyr::filter() masks stats::filter()  
## X dplyr::lag()     masks stats::lag()
```

```
queens_components <- get_estimates(  
  geography = "county",  
  product = "components",  
  state = "NY",  
  county = "Queens",  
  year = 2019  
)
```

```
print(queens_components)
```

```
## # A tibble: 12 × 4  
##   NAME                GEOID variable      value  
##   <chr>              <chr> <chr>      <dbl>  
## 1 Queens County, New York 36081 BIRTHS      27453  
## 2 Queens County, New York 36081 DEATHS     16380  
## 3 Queens County, New York 36081 DOMESTICMIG -41789  
## 4 Queens County, New York 36081 INTERNATIONALMIG  9883  
## 5 Queens County, New York 36081 NATURALINC     11073  
## 6 Queens County, New York 36081 NETMIG     -31906  
## 7 Queens County, New York 36081 RBIRTH        12.1  
## 8 Queens County, New York 36081 RDEATH         7.23  
## 9 Queens County, New York 36081 RDOMESTICMIG   -18.5  
## 10 Queens County, New York 36081 RINTERNATIONALMIG  4.36  
## 11 Queens County, New York 36081 RNATURALINC     4.89  
## 12 Queens County, New York 36081 RNETMIG      -14.1
```

```
louisiana_sex_hisp <- get_estimates(  
  geography = "state",  
  product = "characteristics",  
  breakdown = c("SEX", "HISP"),  
  breakdown_labels = TRUE,  
  state = "LA",  
  year = 2019  
)
```

```
print(louisiana_sex_hisp)
```

```
## # A tibble: 9 × 5
##   GEOID NAME      value SEX      HISP
##   <chr> <chr>      <dbl> <chr>   <chr>
## 1 22 Louisiana 4648794 Both sexes Both Hispanic Origins
## 2 22 Louisiana 4401822 Both sexes Non-Hispanic
## 3 22 Louisiana 246972 Both sexes Hispanic
## 4 22 Louisiana 2267050 Male      Both Hispanic Origins
## 5 22 Louisiana 2135979 Male      Non-Hispanic
## 6 22 Louisiana 131071 Male      Hispanic
## 7 22 Louisiana 2381744 Female    Both Hispanic Origins
## 8 22 Louisiana 2265843 Female    Non-Hispanic
## 9 22 Louisiana 115901 Female    Hispanic
```

```
honolulu_migration <- get_flows(
  geography = "county",
  state = "HI",
  county = "Honolulu",
  year = 2019
)

print(honolulu_migration)
```

```
## # A tibble: 3,156 × 7
##   GEOID1 GEOID2 FULL1_NAME      FULL2_NAME      variable estimate moe
##   <chr> <chr> <chr>      <chr>      <chr>      <dbl> <dbl>
## 1 15003 <NA> Honolulu County, Hawaii Africa      MOVEDIN      152 156
## 2 15003 <NA> Honolulu County, Hawaii Africa      MOVEDOUT      NA  NA
## 3 15003 <NA> Honolulu County, Hawaii Africa      MOVEDNET      NA  NA
## 4 15003 <NA> Honolulu County, Hawaii Asia      MOVEDIN     7680 884
## 5 15003 <NA> Honolulu County, Hawaii Asia      MOVEDOUT      NA  NA
## 6 15003 <NA> Honolulu County, Hawaii Asia      MOVEDNET      NA  NA
## 7 15003 <NA> Honolulu County, Hawaii Central America MOVEDIN      192 100
## 8 15003 <NA> Honolulu County, Hawaii Central America MOVEDOUT      NA  NA
## 9 15003 <NA> Honolulu County, Hawaii Central America MOVEDNET      NA  NA
## 10 15003 <NA> Honolulu County, Hawaii Caribbean      MOVEDIN       97 78
## # ... with 3,146 more rows
```

```
cbsa_bachelors <- get_acs(
  geography = "cbsa",
  variables = "DP02_0068P",
  year = 2019,
  show_call = TRUE
)
```

```
## Getting data from the 2015-2019 5-year ACS
```

```
## Using the ACS Data Profile
```

```
## Census API call: https://api.census.gov/data/2019/acs/acs5/profile?get=DP02_0068PE%2CDP02_0068PM%2CNAME&for=metropolitan%20statistical%20area%2Fmicropolitan%20statistical%20area%3A%2A
```

```
print(cbsa_bachelors)
```

```
## # A tibble: 938 × 5
```

##	GEOID	NAME	variable	estimate	moe
##	<chr>	<chr>	<chr>	<dbl>	<dbl>
## 1	10100	Aberdeen, SD Micro Area	DP02_0068P	29.1	2.1
## 2	10140	Aberdeen, WA Micro Area	DP02_0068P	16.5	1.3
## 3	10180	Abilene, TX Metro Area	DP02_0068P	23	1
## 4	10220	Ada, OK Micro Area	DP02_0068P	28.2	1.5
## 5	10300	Adrian, MI Micro Area	DP02_0068P	21	0.9
## 6	10380	Aguadilla-Isabela, PR Metro Area	DP02_0068P	NA	NA
## 7	10420	Akron, OH Metro Area	DP02_0068P	31.7	0.5
## 8	10460	Alamogordo, NM Micro Area	DP02_0068P	17.5	1.5
## 9	10500	Albany, GA Metro Area	DP02_0068P	20	0.9
## 10	10540	Albany-Lebanon, OR Metro Area	DP02_0068P	19.3	1.2
## #	... with 928 more rows				

Wrangling Census data with tidyverse tools

Ram Mandava

2023-02-17

Exploring Census data with tidyverse tools

```
library(tidycensus)
```

```
## Warning: package 'tidycensus' was built under R version 4.2.2
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.2.2
```

```
## — Attaching packages — tidyverse 1.3.2 —
## ✓ ggplot2 3.4.0      ✓ purrr  1.0.1
## ✓ tibble  3.1.8      ✓ dplyr  1.0.10
## ✓ tidyr   1.3.0      ✓ stringr 1.5.0
## ✓ readr   2.1.3      ✓ forcats 1.0.0
```

```
## Warning: package 'ggplot2' was built under R version 4.2.2
```

```
## Warning: package 'tidyr' was built under R version 4.2.2
```

```
## Warning: package 'readr' was built under R version 4.2.2
```

```
## Warning: package 'purrr' was built under R version 4.2.2
```

```
## Warning: package 'stringr' was built under R version 4.2.2
```

```
## Warning: package 'forcats' was built under R version 4.2.2
```

```
## — Conflicts — tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
```



```
median_age <- get_acs(
  geography = "county",
  variables = "B01002_001",
  year = 2020
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
arrange(median_age, estimate)
```

```
## # A tibble: 3,221 × 5
##   GEOID NAME                variable estimate moe
##   <chr> <chr>                <chr>      <dbl> <dbl>
## 1 35011 De Baca County, New Mexico B01002_001    22.2  6.9
## 2 51678 Lexington city, Virginia  B01002_001    22.2  0.8
## 3 16065 Madison County, Idaho     B01002_001    23.5  0.2
## 4 46121 Todd County, South Dakota B01002_001    23.6  0.6
## 5 51750 Radford city, Virginia    B01002_001    23.7  0.6
## 6 13053 Chattahoochee County, Georgia B01002_001    24    0.7
## 7 02158 Kusilvak Census Area, Alaska B01002_001    24.1  0.2
## 8 49049 Utah County, Utah         B01002_001    25    0.1
## 9 46027 Clay County, South Dakota  B01002_001    25.2  0.5
## 10 53075 Whitman County, Washington B01002_001    25.2  0.2
## # ... with 3,211 more rows
```

```
arrange(median_age, desc(estimate))
```

```
## # A tibble: 3,221 × 5
##   GEOID NAME                variable estimate moe
##   <chr> <chr>                <chr>      <dbl> <dbl>
## 1 12119 Sumter County, Florida     B01002_001    68    0.3
## 2 48301 Loving County, Texas       B01002_001   62.2  37.8
## 3 48243 Jeff Davis County, Texas   B01002_001   61.3  36.8
## 4 08027 Custer County, Colorado    B01002_001   60.1  3.4
## 5 12015 Charlotte County, Florida  B01002_001   59.5  0.2
## 6 51091 Highland County, Virginia  B01002_001   59.5  4.8
## 7 35003 Catron County, New Mexico   B01002_001   59.4  2.6
## 8 51133 Northumberland County, Virginia B01002_001   59.3  0.7
## 9 26131 Ontonagon County, Michigan B01002_001   59.1  0.5
## 10 48443 Terrell County, Texas      B01002_001   59.1  9.4
## # ... with 3,211 more rows
```

```
filter(median_age, estimate >= 50)
```

```
## # A tibble: 218 × 5
##   GEOID NAME                variable estimate moe
##   <chr> <chr>                <chr>      <dbl> <dbl>
## 1 02105 Hoonah-Angoon Census Area, Alaska B01002_001 52.1 2.9
## 2 04007 Gila County, Arizona B01002_001 50.4 0.2
## 3 04012 La Paz County, Arizona B01002_001 57.4 0.6
## 4 04015 Mohave County, Arizona B01002_001 52.3 0.2
## 5 04025 Yavapai County, Arizona B01002_001 54.1 0.2
## 6 05005 Baxter County, Arkansas B01002_001 52.3 0.5
## 7 05089 Marion County, Arkansas B01002_001 52.1 0.8
## 8 05097 Montgomery County, Arkansas B01002_001 50.6 0.8
## 9 05137 Stone County, Arkansas B01002_001 50 0.5
## 10 06009 Calaveras County, California B01002_001 52.8 0.6
## # ... with 208 more rows
```

```
separate(
  median_age,
  NAME,
  into = c("county", "state"),
  sep = ", "
)
```

```
## # A tibble: 3,221 × 6
##   GEOID county      state variable estimate moe
##   <chr> <chr>      <chr> <chr>      <dbl> <dbl>
## 1 01001 Autauga County Alabama B01002_001 38.6 0.6
## 2 01003 Baldwin County Alabama B01002_001 43.2 0.4
## 3 01005 Barbour County Alabama B01002_001 40.1 0.6
## 4 01007 Bibb County Alabama B01002_001 39.9 1.2
## 5 01009 Blount County Alabama B01002_001 41 0.5
## 6 01011 Bullock County Alabama B01002_001 39.7 1.9
## 7 01013 Butler County Alabama B01002_001 41.2 0.6
## 8 01015 Calhoun County Alabama B01002_001 39.5 0.4
## 9 01017 Chambers County Alabama B01002_001 41.9 0.7
## 10 01019 Cherokee County Alabama B01002_001 46.8 0.5
## # ... with 3,211 more rows
```

```

race_vars <- c(
  White = "B03002_003",
  Black = "B03002_004",
  Native = "B03002_005",
  Asian = "B03002_006",
  HIPI = "B03002_007",
  Hispanic = "B03002_012"
)

az_race <- get_acs(
  geography = "county",
  state = "AZ",
  variables = race_vars,
  summary_var = "B03002_001",
  year = 2020
)

```

```
## Getting data from the 2016-2020 5-year ACS
```

```

az_race_percent <- az_race %>%
  mutate(percent = 100 * (estimate / summary_est)) %>%
  select(NAME, variable, percent)

largest_group <- az_race_percent %>%
  group_by(NAME) %>%
  filter(percent == max(percent))

az_race_percent %>%
  group_by(variable) %>%
  summarize(median_pct = median(percent))

```

```

## # A tibble: 6 × 2
##   variable median_pct
##   <chr>         <dbl>
## 1 Asian         0.992
## 2 Black         1.29
## 3 HIPI          0.107
## 4 Hispanic     30.5
## 5 Native        3.63
## 6 White        53.8

```

```

mn_hh_income <- get_acs(
  geography = "county",
  table = "B19001",
  state = "MN",
  year = 2016
)

```

```
## Getting data from the 2012-2016 5-year ACS
```

```
mn_hh_income_recode <- mn_hh_income %>%  
  filter(variable != "B19001_001") %>%  
  mutate(incgroup = case_when(  
    variable < "B19001_008" ~ "below35k",  
    variable < "B19001_013" ~ "bw35kand75k",  
    TRUE ~ "above75k"  
  ))  
  
mn_group_sums <- mn_hh_income_recode %>%  
  group_by(GEOID, incgroup) %>%  
  summarize(estimate = sum(estimate))
```

```
## `summarise()` has grouped output by 'GEOID'. You can override using the  
## `.groups` argument.
```

```
oglala_lakota_age <- get_acs(  
  geography = "county",  
  state = "SD",  
  county = "Oglala Lakota",  
  table = "B01001",  
  year = 2020  
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
## Loading ACS5 variables for 2020 from table B01001. To cache this dataset for faster access to  
ACS tables in the future, run this function with `cache_table = TRUE`. You only need to do this  
once per ACS dataset.
```

```
oglala_lakota_age_10 <- get_acs(  
  geography = "county",  
  state = "SD",  
  county = "Shannon",  
  table = "B01001",  
  year = 2010  
)
```

```
## Getting data from the 2006-2010 5-year ACS
```

```
## Loading ACS5 variables for 2010 from table B01001. To cache this dataset for faster access to  
ACS tables in the future, run this function with `cache_table = TRUE`. You only need to do this  
once per ACS dataset.
```

```
co_college19 <- get_acs(  
  geography = "county",  
  variables = "DP02_0068P",  
  state = "CO",  
  survey = "acs1",  
  year = 2019  
)
```

```
## Getting data from the 2019 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Using the ACS Data Profile
```

```
co_college18 <- get_acs(  
  geography = "county",  
  variables = "DP02_0068P",  
  state = "CO",  
  survey = "acs1",  
  year = 2018  
)
```

```
## Getting data from the 2018 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Using the ACS Data Profile
```

```
ak_income_compare <- get_acs(  
  geography = "county",  
  variables = c(  
    income15 = "CP03_2015_062",  
    income20 = "CP03_2020_062"  
  ),  
  state = "AK",  
  year = 2020  
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
## Using the ACS Comparison Profile
```

```
college_vars <- c("B15002_015",
                  "B15002_016",
                  "B15002_017",
                  "B15002_018",
                  "B15002_032",
                  "B15002_033",
                  "B15002_034",
                  "B15002_035")

years <- 2010:2019
names(years) <- years

college_by_year <- map_dfr(years, ~{
  get_acs(
    geography = "county",
    variables = college_vars,
    state = "CO",
    summary_var = "B15002_001",
    survey = "acs1",
    year = .x
  )
}, .id = "year")
```

```
## Getting data from the 2010 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2011 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2012 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2013 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2014 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2015 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2016 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2017 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2018 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2019 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
college_by_year %>%  
  arrange(NAME, variable, year)
```

```
## # A tibble: 920 × 8  
##   year GEOID NAME          variable estimate moe summar...1 summa...2  
##   <chr> <chr> <chr>          <chr>      <dbl> <dbl>    <dbl>    <dbl>  
## 1 2010 08001 Adams County, Colorado B15002_015 20501 1983 275849 790  
## 2 2011 08001 Adams County, Colorado B15002_015 21233 2124 281231 865  
## 3 2012 08001 Adams County, Colorado B15002_015 19238 2020 287924 693  
## 4 2013 08001 Adams County, Colorado B15002_015 23818 2445 295122 673  
## 5 2014 08001 Adams County, Colorado B15002_015 20255 1928 304394 541  
## 6 2015 08001 Adams County, Colorado B15002_015 22962 2018 312281 705  
## 7 2016 08001 Adams County, Colorado B15002_015 25744 2149 318077 525  
## 8 2017 08001 Adams County, Colorado B15002_015 26159 2320 324185 562  
## 9 2018 08001 Adams County, Colorado B15002_015 28113 2078 331247 955  
## 10 2019 08001 Adams County, Colorado B15002_015 27552 2070 336931 705  
## # ... with 910 more rows, and abbreviated variable names 1summary_est,  
## # 2summary_moe
```

```
percent_college_by_year <- college_by_year %>%
  group_by(NAME, year) %>%
  summarize(numerator = sum(estimate),
            denominator = first(summary_est)) %>%
  mutate(pct_college = 100 * (numerator / denominator)) %>%
  pivot_wider(id_cols = NAME,
              names_from = year,
              values_from = pct_college)
```

```
## `summarise()` has grouped output by 'NAME'. You can override using the
## `.groups` argument.
```

```
get_acs(
  geography = "county",
  state = "Rhode Island",
  variables = "B19013_001",
  year = 2020
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
## # A tibble: 5 × 5
##   GEOID NAME                variable estimate moe
##   <chr> <chr>                <chr>      <dbl> <dbl>
## 1 44001 Bristol County, Rhode Island B19013_001 85413 6122
## 2 44003 Kent County, Rhode Island    B19013_001 75857 2022
## 3 44005 Newport County, Rhode Island B19013_001 84282 2629
## 4 44007 Providence County, Rhode Island B19013_001 62323 1270
## 5 44009 Washington County, Rhode Island B19013_001 86970 3651
```

```
get_acs(
  geography = "county",
  state = "Rhode Island",
  variables = "B19013_001",
  year = 2020,
  moe_level = 99
)
```

```
## Getting data from the 2016-2020 5-year ACS
```



```
## # A tibble: 5 × 5
##   GEOID NAME                variable estimate   moe
##   <chr> <chr>                <chr>      <dbl> <dbl>
## 1 44001 Bristol County, Rhode Island B19013_001 85413 9527.
## 2 44003 Kent County, Rhode Island    B19013_001 75857 3147.
## 3 44005 Newport County, Rhode Island B19013_001 84282 4091.
## 4 44007 Providence County, Rhode Island B19013_001 62323 1976.
## 5 44009 Washington County, Rhode Island B19013_001 86970 5682.
```

```
vars <- paste0("B01001_0", c(20:25, 44:49))
```

```
vars
```

```
## [1] "B01001_020" "B01001_021" "B01001_022" "B01001_023" "B01001_024"
## [6] "B01001_025" "B01001_044" "B01001_045" "B01001_046" "B01001_047"
## [11] "B01001_048" "B01001_049"
```

```
salt_lake <- get_acs(
  geography = "tract",
  variables = vars,
  state = "Utah",
  county = "Salt Lake",
  year = 2020
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
example_tract <- salt_lake %>%
  filter(GEOID == "49035100100")

example_tract %>%
  select(-NAME)
```

```
## # A tibble: 12 × 4
##   GEOID      variable estimate   moe
##   <chr>      <chr>      <dbl> <dbl>
## 1 49035100100 B01001_020         11    13
## 2 49035100100 B01001_021         25    18
## 3 49035100100 B01001_022          7    10
## 4 49035100100 B01001_023          4     7
## 5 49035100100 B01001_024          0    12
## 6 49035100100 B01001_025         17    20
## 7 49035100100 B01001_044          0    12
## 8 49035100100 B01001_045          4     7
## 9 49035100100 B01001_046         21    17
## 10 49035100100 B01001_047        123   168
## 11 49035100100 B01001_048         17    21
## 12 49035100100 B01001_049          0    12
```

```
moe_prop(25, 100, 5, 3)
```

```
## [1] 0.0494343
```

```
salt_lake_grouped <- salt_lake %>%
  mutate(sex = case_when(
    str_sub(variable, start = -2) < "26" ~ "Male",
    TRUE ~ "Female"
  )) %>%
  group_by(GEOID, sex) %>%
  summarize(sum_est = sum(estimate),
            sum_moe = moe_sum(moe, estimate))
```

```
## `summarise()` has grouped output by 'GEOID'. You can override using the
## `.groups` argument.
```

Exploring US Census data with visualization

Ram Mandava

2023-02-17

Basic Census visualization with ggplot2

```
library(tidycensus)
```

```
## Warning: package 'tidycensus' was built under R version 4.2.2
```

```
ga_wide <- get_acs(  
  geography = "county",  
  state = "Georgia",  
  variables = c(medinc = "B19013_001",  
                medage = "B01002_001"),  
  output = "wide",  
  year = 2020  
)
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.2.2
```

```
## — Attaching packages — tidyverse 1.3.2 —  
## ✓ ggplot2 3.4.0      ✓ purrr   1.0.1  
## ✓ tibble  3.1.8      ✓ dplyr   1.0.10  
## ✓ tidyr   1.3.0      ✓ stringr 1.5.0  
## ✓ readr   2.1.3      ✓ forcats 1.0.0
```

```
## Warning: package 'ggplot2' was built under R version 4.2.2
```

```
## Warning: package 'tidyr' was built under R version 4.2.2
```

```
## Warning: package 'readr' was built under R version 4.2.2
```

```
## Warning: package 'purrr' was built under R version 4.2.2
```

```
## Warning: package 'stringr' was built under R version 4.2.2
```


```
## Warning: package 'forcats' was built under R version 4.2.2
```

```
## — Conflicts ————— tidyverse_conflicts() —  
## X dplyr::filter() masks stats::filter()  
## X dplyr::lag()    masks stats::lag()
```


```
options(scipen = 999)
```

```
ggplot(ga_wide, aes(x = medincE)) +  
  geom_histogram()
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
ggplot(ga_wide, aes(x = medincE)) +  
  geom_histogram(bins = 15)
```



```
ggplot(ga_wide, aes(y = medincE)) +  
  geom_boxplot()
```




```
ggplot(ga_wide, aes(x = medageE, y = medincE)) +  
  geom_point()
```



```
ggplot(ga_wide, aes(x = medageE, y = medincE)) +  
  geom_point() +  
  geom_smooth(method = "lm")
```

```
## `geom_smooth()` using formula = 'y ~ x'
```



```
library(tidycensus)
library(tidyverse)


metros <- get_acs(
  geography = "cbsa",
  variables = "DP03_0021P",
  summary_var = "B01003_001",
  survey = "acs1",
  year = 2019
) %>%
  slice_max(summary_est, n = 20)
```

```
## Getting data from the 2019 1-year ACS
```


```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Using the ACS Data Profile
```

```
ggplot(metros, aes(x = NAME, y = estimate)) +
  geom_col()
```

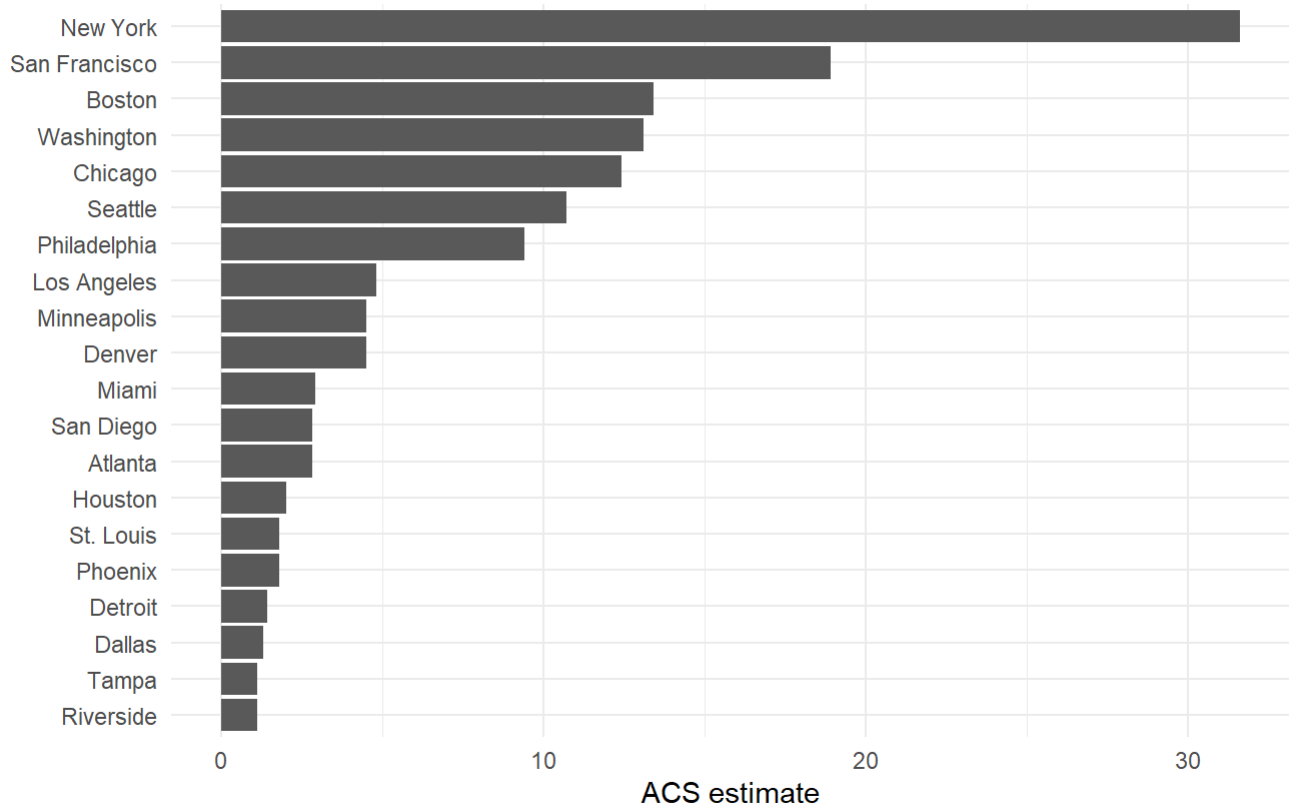



```
metros %>%
  mutate(NAME = str_remove(NAME, "-.*$")) %>%
  mutate(NAME = str_remove(NAME, ",.*$")) %>%
  ggplot(aes(y = reorder(NAME, estimate), x = estimate)) +
  geom_col()
```



```
metros %>%
  mutate(NAME = str_remove(NAME, "-.*$")) %>%
  mutate(NAME = str_remove(NAME, ",.*$")) %>%
  ggplot(aes(y = reorder(NAME, estimate), x = estimate)) +
  geom_col() +
  theme_minimal() +
  labs(title = "Public transit commute share",
       subtitle = "2019 1-year ACS estimates",
       y = "",
       x = "ACS estimate",
       caption = "Source: ACS Data Profile variable DP03_0021P via the tidycensus R package")
```

Public transit commute share 2019 1-year ACS estimates



Source: ACS Data Profile variable DP03_0021P via the tidycensus R package

```
library(scales)
```

```
## Warning: package 'scales' was built under R version 4.2.2
```

```
##  
## Attaching package: 'scales'
```

```
## The following object is masked from 'package:purrr':  
##  
##   discard
```

```
## The following object is masked from 'package:readr':  
##  
##   col_factor
```

```
metros %>%
  mutate(NAME = str_remove(NAME, "-.*$")) %>%
  mutate(NAME = str_remove(NAME, ",.*$")) %>%
  ggplot(aes(y = reorder(NAME, estimate), x = estimate)) +
  geom_col(color = "navy", fill = "navy",
    alpha = 0.5, width = 0.85) +
  theme_minimal(base_size = 12, base_family = "Verdana") +
  scale_x_continuous(labels = label_percent(scale = 1)) +
  labs(title = "Public transit commute share",
    subtitle = "2019 1-year ACS estimates",
    y = "",
    x = "ACS estimate",
    caption = "Source: ACS Data Profile variable DP03_0021P via the tidycensus R package")
```

```
## Warning in grid.Call(C_stringMetric, as.graphicsAnnot(x$label)): font family
## not found in Windows font database
```

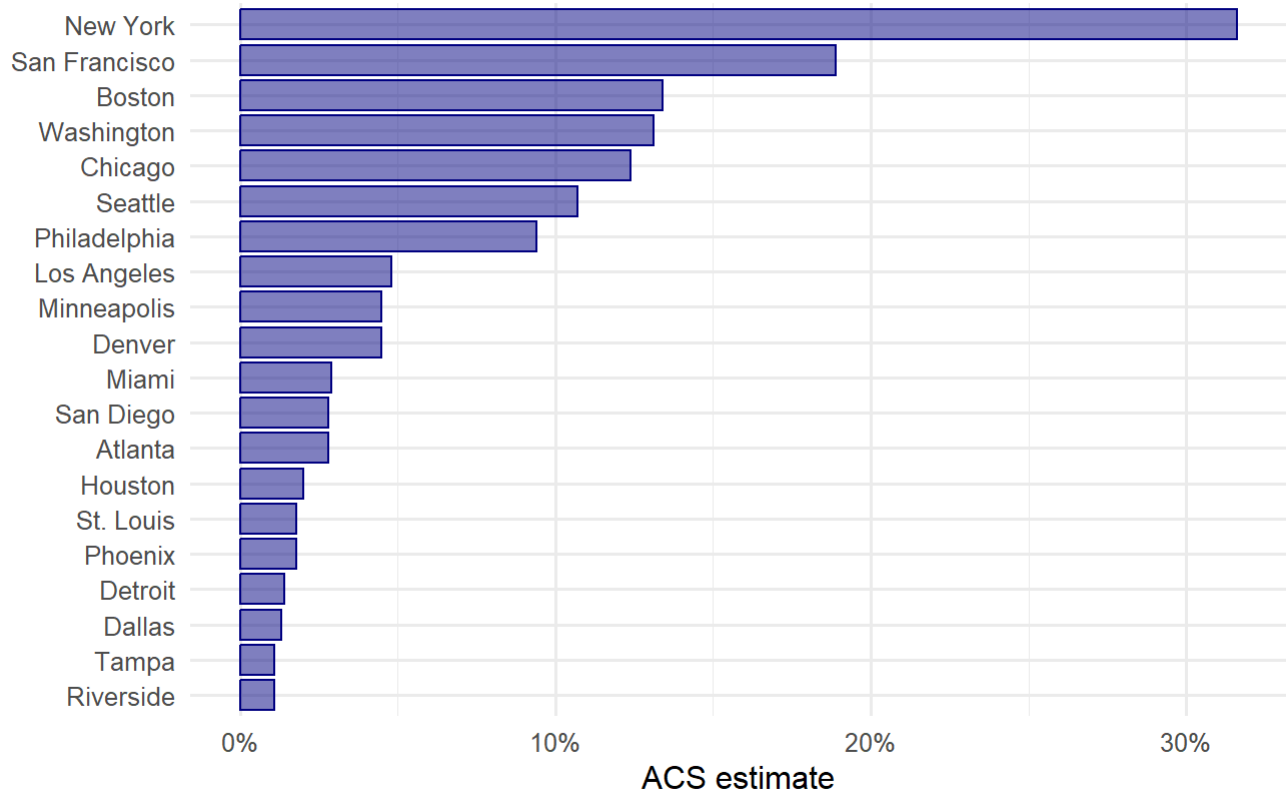
```
## Warning in grid.Call(C_stringMetric, as.graphicsAnnot(x$label)): font family
## not found in Windows font database
```

```
## Warning in grid.Call(C_stringMetric, as.graphicsAnnot(x$label)): font family
## not found in Windows font database
```

[illegible]

Public transit commute share

2019 1-year ACS estimates



Source: ACS Data Profile variable DP03_0021P via the tidycensus R package

```
ggsave("metro_transit.png")
```

```
## Saving 7 x 5 in image
```

[illegible]

```
ggsave(  
  filename = "metro_transit.png",  
  path = "~/images",  
  width = 8,  
  height = 5,  
  units = "in",  
  dpi = 300  
)
```


[illegible]

```
maine <- get_decennial(  
  state = "Maine",  
  geography = "county",  
  variables = c(totalpop = "P1_001N"),  
  year = 2020  
) %>%  
  arrange(desc(value))
```

```
## Getting data from the 2020 decennial Census
```

```
## Using the PL 94-171 Redistricting Data summary file
```

```
## Note: 2020 decennial Census data use differential privacy, a technique that  
## introduces errors into data to preserve respondent confidentiality.  
## i Small counts should be interpreted with caution.  
## i See https://www.census.gov/library/fact-sheets/2021/protecting-the-confidentiality-of-the-2020-census-redistricting-data.html for additional guidance.  
## This message is displayed once per session.
```


```
maine_income <- get_acs(  
  state = "Maine",  
  geography = "county",  
  variables = c(hhincome = "B19013_001"),  
  year = 2020  
) %>%  
  mutate(NAME = str_remove(NAME, " County, Maine"))
```

```
## Getting data from the 2016-2020 5-year ACS
```

```
ggplot(maine_income, aes(x = estimate, y = reorder(NAME, estimate))) +  
  geom_point(size = 3, color = "darkgreen") +  
  labs(title = "Median household income",  
       subtitle = "Counties in Maine",  
       x = "",  
       y = "ACS estimate") +  
  theme_minimal(base_size = 12.5) +  
  scale_x_continuous(labels = label_dollar())
```

Median household income

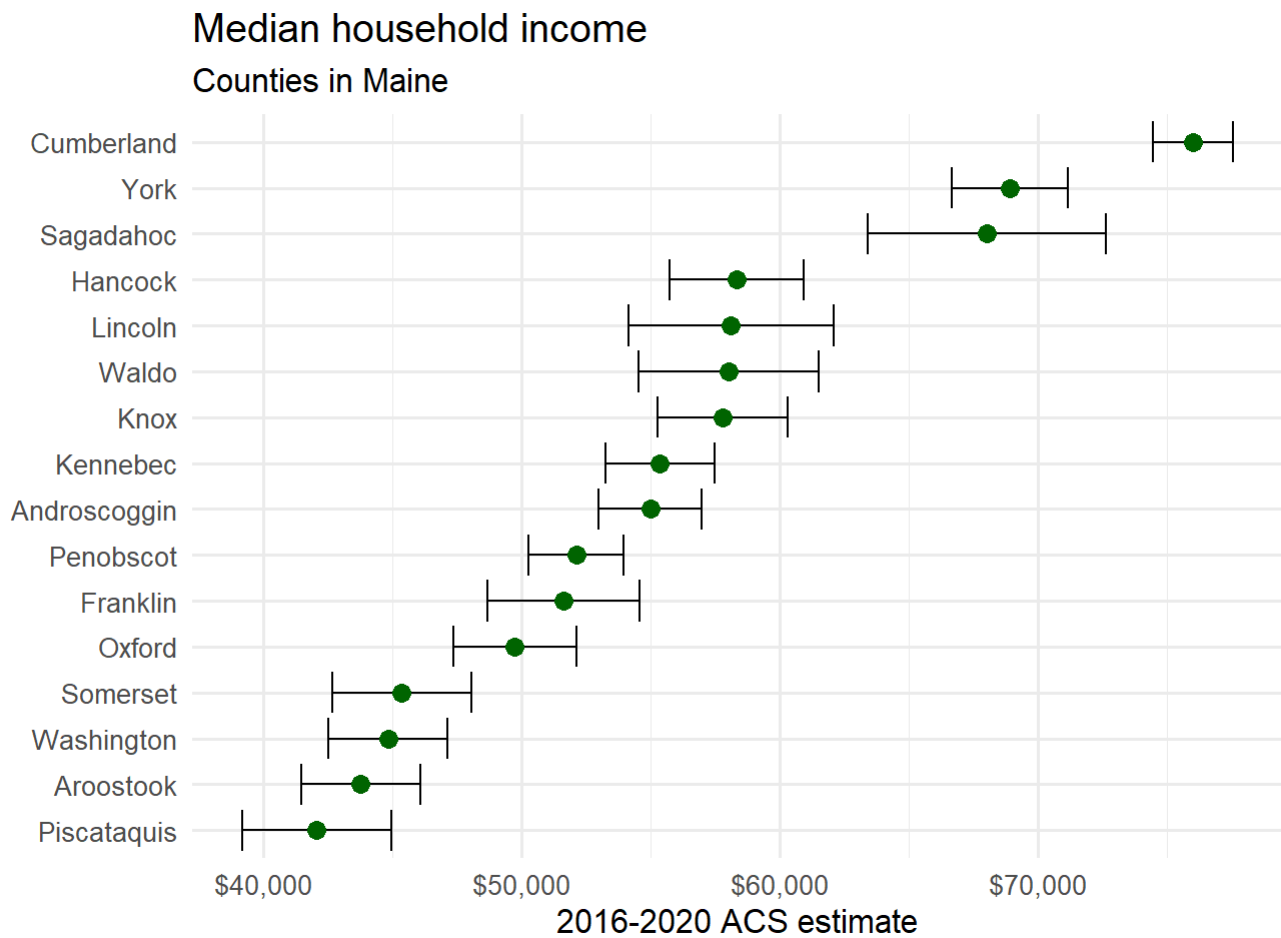
Counties in Maine



```
maine_income %>%
  arrange(desc(moe))
```

```
## # A tibble: 16 × 5
##   GEOID NAME      variable estimate   moe
##   <chr> <chr>      <chr>      <dbl> <dbl>
## 1 23023 Sagadahoc hhincome    68039  4616
## 2 23015 Lincoln   hhincome    58125  3974
## 3 23027 Waldo     hhincome    58034  3482
## 4 23007 Franklin  hhincome    51630  2948
## 5 23021 Piscataquis hhincome    42083  2883
## 6 23025 Somerset  hhincome    45382  2694
## 7 23009 Hancock   hhincome    58345  2593
## 8 23013 Knox      hhincome    57794  2528
## 9 23017 Oxford    hhincome    49761  2380
## 10 23003 Aroostook hhincome    43791  2306
## 11 23029 Washington hhincome    44847  2292
## 12 23031 York      hhincome    68932  2239
## 13 23011 Kennebec  hhincome    55368  2112
## 14 23001 Androscoggin hhincome    55002  2003
## 15 23019 Penobscot hhincome    52128  1836
## 16 23005 Cumberland hhincome    76014  1563
```

```
ggplot(maine_income, aes(x = estimate, y = reorder(NAME, estimate))) +
  geom_errorbarh(aes(xmin = estimate - moe, xmax = estimate + moe)) +
  geom_point(size = 3, color = "darkgreen") +
  theme_minimal(base_size = 12.5) +
  labs(title = "Median household income",
       subtitle = "Counties in Maine",
       x = "2016-2020 ACS estimate",
       y = "") +
  scale_x_continuous(labels = label_dollar())
```



```
years <- 2005:2019
names(years) <- years

deschutes_value <- map_dfr(years, ~{
  get_acs(
    geography = "county",
    variables = "B25077_001",
    state = "OR",
    county = "Deschutes",
    year = .x,
    survey = "acs1"
  )
}, .id = "year")
```

Getting data from the 2005 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2006 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2007 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2008 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2009 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2010 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2011 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2012 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2013 1-year ACS

The 1-year ACS provides data for geographies with populations of 65,000 and greater.

Getting data from the 2014 1-year ACS

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2015 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2016 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2017 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```


```
## Getting data from the 2018 1-year ACS
```

```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
## Getting data from the 2019 1-year ACS
```


```
## The 1-year ACS provides data for geographies with populations of 65,000 and greater.
```

```
ggplot(deschutes_value, aes(x = year, y = estimate, group = 1)) +  
  geom_line() +  
  geom_point()
```




```
ggplot(deschutes_value, aes(x = year, y = estimate, group = 1)) +
  geom_ribbon(aes(ymax = estimate + moe, ymin = estimate - moe),
            fill = "navy",
            alpha = 0.4) +
  geom_line(color = "navy") +
  geom_point(color = "navy", size = 2) +
  theme_minimal(base_size = 12) +
  scale_y_continuous(labels = label_dollar(scale = .001, suffix = "k")) +
  labs(title = "Median home value in Deschutes County, OR",
       x = "Year",
       y = "ACS estimate",
       caption = "Shaded area represents margin of error around the ACS estimate")
```

Median home value in Deschutes County, OR



Shaded area represents margin of error around the ACS estimate

```
utah <- get_estimates(  
  geography = "state",  
  state = "UT",  
  product = "characteristics",  
  breakdown = c("SEX", "AGEGROUP"),  
  breakdown_labels = TRUE,  
  year = 2019  
)  
  
utah_filtered <- filter(utah, str_detect(AGEGROUP, "^Age"),  
  SEX != "Both sexes") %>%  
  mutate(value = ifelse(SEX == "Male", -value, value))  
  
ggplot(utah_filtered, aes(x = value, y = AGEGROUP, fill = SEX)) +  
  geom_col()
```

```

utah_pyramid <- ggplot(utah_filtered,
  aes(x = value,
    y = AGEGROUP,
    fill = SEX)) +
  geom_col(width = 0.95, alpha = 0.75) +
  theme_minimal(base_family = "Verdana",
    base_size = 12) +
  scale_x_continuous(
    labels = ~ number_format(scale = .001, suffix = "k")(abs(.x)),
    limits = 140000 * c(-1,1)
  ) +
  scale_y_discrete(labels = ~ str_remove_all(.x, "Age\\s|\\syears")) +
  scale_fill_manual(values = c("darkred", "navy")) +
  labs(x = "",
    y = "2019 Census Bureau population estimate",
    title = "Population structure in Utah",
    fill = "",
    caption = "Data source: US Census Bureau population estimates & tidycensus R package")

utah_pyramid

```




```
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, : font
## family not found in Windows font database

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, : font
## family not found in Windows font database

## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, : font
## family not found in Windows font database
```

Population structure in Utah



Data source: US Census Bureau population estimates & tidycensus R package

```
housing_val <- get_acs(
  geography = "tract",
  variables = "B25077_001",
  state = "OR",
  county = c(
    "Multnomah",
    "Clackamas",
    "Washington",
    "Yamhill",
    "Marion",
    "Columbia"
  ),
  year = 2020
)
```


```
## Getting data from the 2016-2020 5-year ACS
```

```
housing_val2 <- separate(  
  housing_val,  
  NAME,  
  into = c("tract", "county", "state"),  
  sep = ", "  
)  
  
housing_val2 %>%  
  group_by(county) %>%  
  summarize(min = min(estimate, na.rm = TRUE),  
            mean = mean(estimate, na.rm = TRUE),  
            median = median(estimate, na.rm = TRUE),  
            max = max(estimate, na.rm = TRUE))
```

```
## # A tibble: 6 × 5  
##   county          min    mean median    max  
##   <chr>          <dbl>   <dbl> <dbl>   <dbl>  
## 1 Clackamas County 62800 449941. 426700 909000  
## 2 Columbia County 218100 277591. 275900 362200  
## 3 Marion County   48700 270969. 261200 483500  
## 4 Multnomah County 192900 455706. 425950 1033500  
## 5 Washington County 221900 419618. 406100 769700  
## 6 Yamhill County   230000 333816. 291100 545500
```

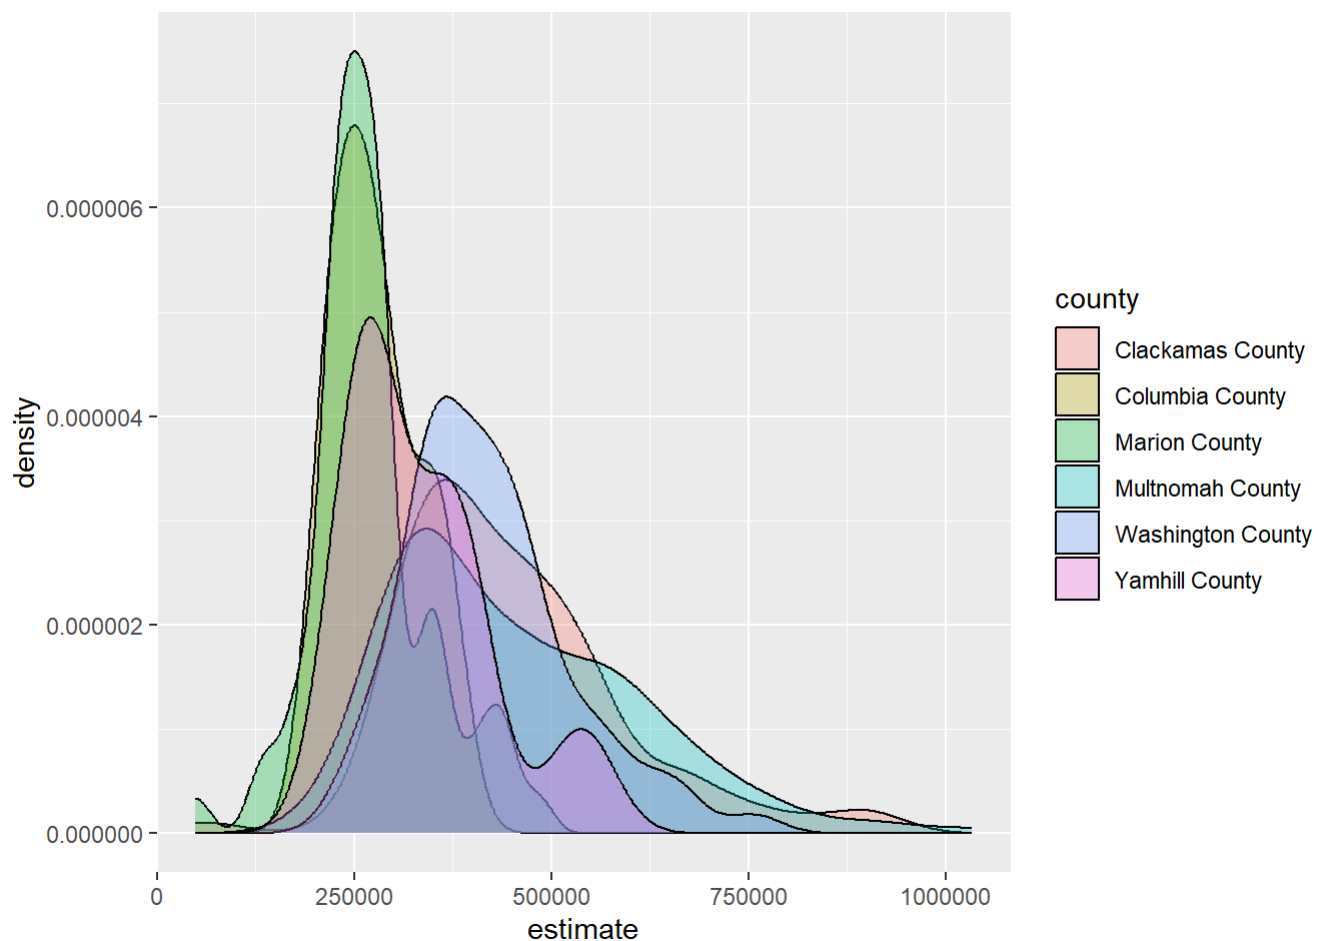
```
ggplot(housing_val2, aes(x = estimate)) +  
  geom_density()
```

```
## Warning: Removed 9 rows containing non-finite values (`stat_density()`).
```



```
ggplot(housing_val2, aes(x = estimate, fill = county)) +  
  geom_density(alpha = 0.3)
```

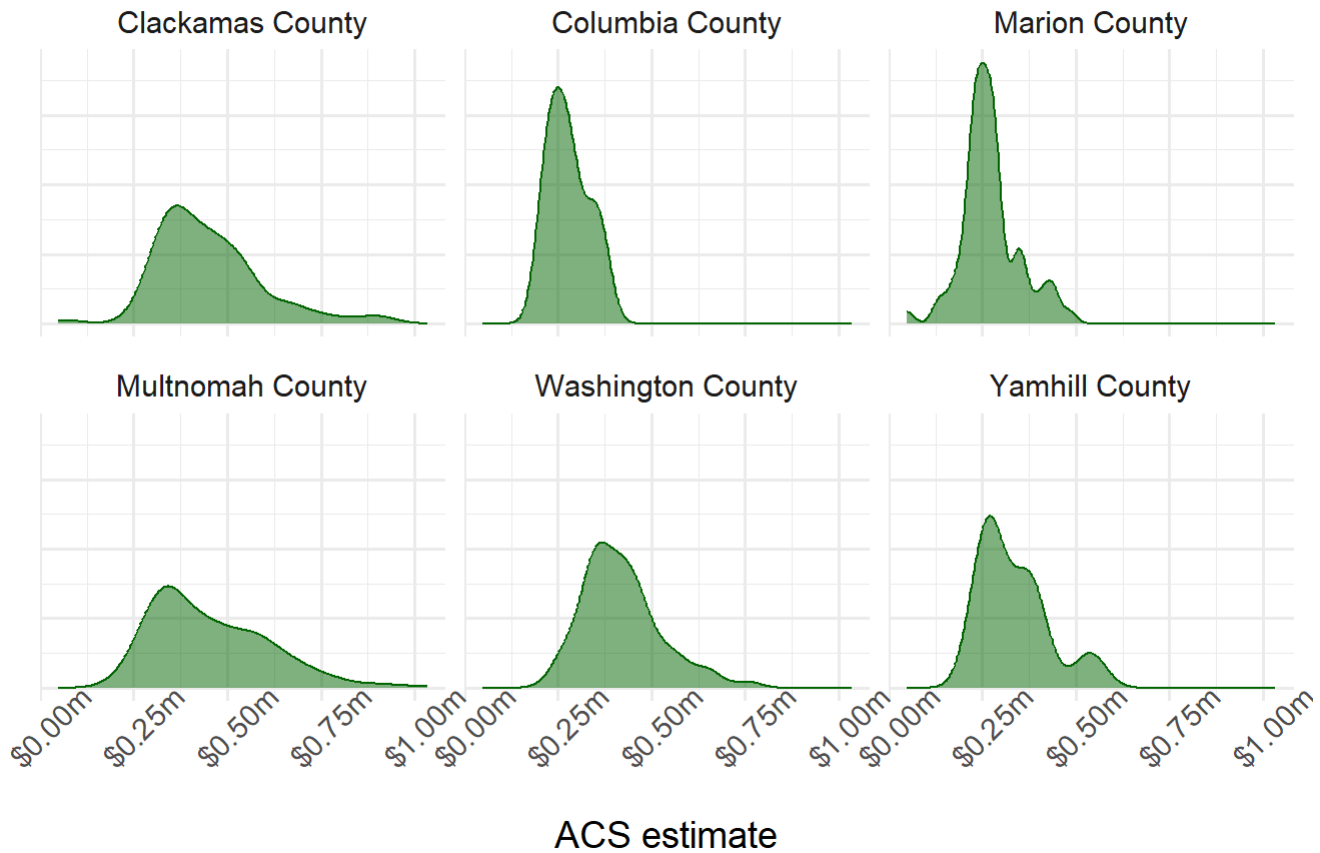
```
## Warning: Removed 9 rows containing non-finite values (`stat_density()`).
```



```
ggplot(housing_val2, aes(x = estimate)) +
  geom_density(fill = "darkgreen", color = "darkgreen", alpha = 0.5) +
  facet_wrap(~county) +
  scale_x_continuous(labels = dollar_format(scale = 0.000001,
                                             suffix = "m")) +
  theme_minimal(base_size = 14) +
  theme(axis.text.y = element_blank(),
        axis.text.x = element_text(angle = 45)) +
  labs(x = "ACS estimate",
       y = "",
       title = "Median home values by Census tract, 2015-2019 ACS")
```

```
## Warning: Removed 9 rows containing non-finite values (`stat_density()`).
```

Median home values by Census tract, 2015-2019 ACS



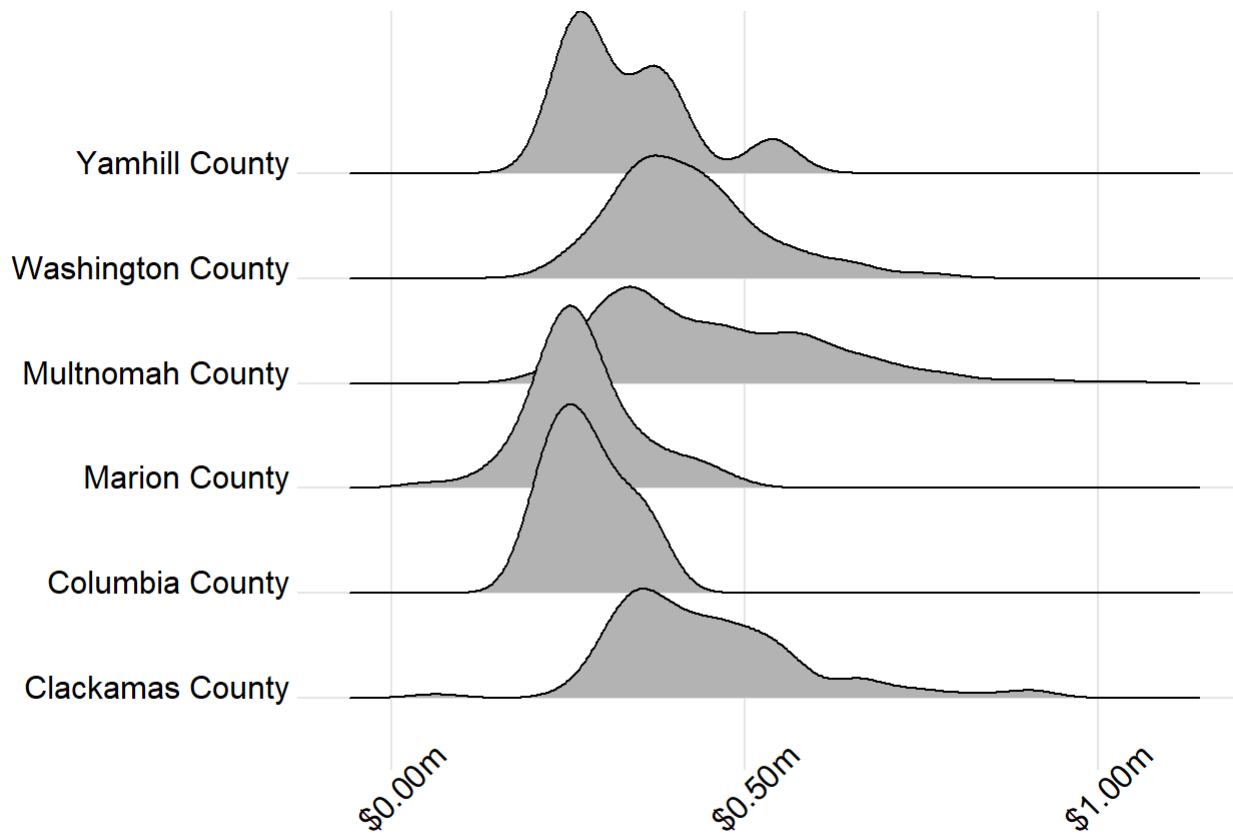
```
library(ggribes)
```

```
## Warning: package 'ggribes' was built under R version 4.2.2
```

```
ggplot(housing_val2, aes(x = estimate, y = county)) +  
  geom_density_ridges() +  
  theme_ridges() +  
  labs(x = "Median home value: 2016-2020 ACS estimate",  
       y = "") +  
  scale_x_continuous(labels = label_dollar(scale = .000001, suffix = "m"),  
                    breaks = c(0, 500000, 1000000)) +  
  theme(axis.text.x = element_text(angle = 45))
```

```
## Picking joint bandwidth of 36200
```

```
## Warning: Removed 9 rows containing non-finite values ( `stat_density_ridges()` ).
```



Median home value: 2016-2020 ACS estimate

```
library(ggbeeswarm)
```

```
## Warning: package 'ggbeeswarm' was built under R version 4.2.2
```


```
ny_race_income <- get_acs(
  geography = "tract",
  state = "NY",
  county = c("New York", "Bronx", "Queens", "Richmond", "Kings"),
  variables = c(White = "B03002_003",
                Black = "B03002_004",
                Asian = "B03002_006",
                Hispanic = "B03002_012"),
  summary_var = "B19013_001",
  year = 2020
) %>%
  group_by(GEOID) %>%
  filter(estimate == max(estimate, na.rm = TRUE)) %>%
  ungroup() %>%
  filter(estimate != 0)
```

```
## Getting data from the 2016-2020 5-year ACS
```



```
ggplot(ny_race_income, aes(x = variable, y = summary_est, color = summary_est)) +
  geom_quasirandom(alpha = 0.5) +
  coord_flip() +
  theme_minimal(base_size = 13) +
  scale_color_viridis_c(guide = "none") +
  scale_y_continuous(labels = label_dollar()) +
  labs(x = "Largest group in Census tract",
       y = "Median household income",
       title = "Household income distribution by largest racial/ethnic group",
       subtitle = "Census tracts, New York City",
       caption = "Data source: 2016-2020 ACS")
```

```
## Warning: Removed 35 rows containing missing values.
```



```
library(geofacet)
```

```
## Warning: package 'geofacet' was built under R version 4.2.2
```


```

us_pyramid_data <- get_estimates(
  geography = "state",
  product = "characteristics",
  breakdown = c("SEX", "AGEGROUP"),
  breakdown_labels = TRUE,
  year = 2019
) %>%
  filter(str_detect(AGEGROUP, "^Age"),
         SEX != "Both sexes") %>%
  group_by(NAME) %>%
  mutate(prop = value / sum(value, na.rm = TRUE)) %>%
  ungroup() %>%
  mutate(prop = ifelse(SEX == "Male", -prop, prop))

ggplot(us_pyramid_data, aes(x = prop, y = AGEGROUP, fill = SEX)) +
  geom_col(width = 1) +
  theme_minimal() +
  scale_fill_manual(values = c("darkred", "navy")) +
  facet_geo(~NAME, grid = "us_state_with_DC_PR_grid2",
            label = "code") +
  theme(axis.text = element_blank(),
        strip.text.x = element_text(size = 8)) +
  labs(x = "",
       y = "",
       title = "Population structure by age and sex",
       fill = "",
       caption = "Data source: US Census Bureau population estimates & tidycensus R package")

```

Population structure by age and sex



Data source: US Census Bureau population estimates & tidycensus R package

```
library(plotly)
```

```
## Warning: package 'plotly' was built under R version 4.2.2
```

```
##  
## Attaching package: 'plotly'
```

```
## The following object is masked from 'package:ggplot2':  
##  
## last_plot
```

```
## The following object is masked from 'package:stats':  
##  
## filter
```

```
## The following object is masked from 'package:graphics':  
##  
## layout
```

```
ggplotly(utah_pyramid)
```

Population structure in Utah

Population structure in Utan

